




Visão geral — média $\pm$ desvio padrão por métrica




Distribuição por run (30 execuções) — Random Forest vs Logistic Regression



Random Forest — Bug-inducing vs Clean




Logistic Regression — Bug-inducing vs Clean




Comparação por classe — Random Forest vs Logistic Regression

Classe: bug-inducing











Classe: clean




Random Forest Logistic Regression

Diferença pareada por seed: Random Forest – Logistic Regression
(barras verdes = RF melhor nesse seed; vermelhas = LR melhor)




Comparação geral (radar) — médias das 6 métricas

Random Forest
Logistic Regression




Commit-message quality vs predicted bug outcome
Defect-inducing commits only; bars show mean across seeds $\pm$ std of seed means


Random Forest — defect-inducing commits



Logistic Regression — defect-inducing commits



Random Forest



Logistic Regression

